

LSTM vs Bi-LSTM Report

1. Performance Comparison

Model	Validation Accuracy	Validation Loss	Test Accuracy
LSTM	0.7374	0.5773	0.7097
Bi-LSTM	0.7334	0.5169	0.8046

The performance metrics above reflect the validation and test results from training both models. Although the Bi-LSTM model had similar validation accuracy to the LSTM, it achieved significantly better test accuracy.

```

Training LSTM model...
Epoch 1/5
/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/embedding.py:90: UserWarning: Argument `input_length` is deprecated. Just remove it.
  warnings.warn(
125/125 ----- 16s 98ms/step - accuracy: 0.6878 - loss: 0.6277 - val_accuracy: 0.7374 - val_loss: 0.5773
Epoch 2/5
125/125 ----- 20s 97ms/step - accuracy: 0.6987 - loss: 0.6182 - val_accuracy: 0.7374 - val_loss: 0.5870
Epoch 3/5
125/125 ----- 21s 98ms/step - accuracy: 0.6994 - loss: 0.6139 - val_accuracy: 0.7374 - val_loss: 0.5863
Epoch 4/5
125/125 ----- 20s 91ms/step - accuracy: 0.7001 - loss: 0.6121 - val_accuracy: 0.7374 - val_loss: 0.5803
Epoch 5/5
125/125 ----- 21s 94ms/step - accuracy: 0.7002 - loss: 0.6138 - val_accuracy: 0.7374 - val_loss: 0.5797
105/105 ----- 1s 19ms/step

LSTM Test Accuracy: 0.7097
precision    recall  f1-score   support

   0         0.71         1.00         0.83        2364
   1         0.00         0.00         0.00         967

 accuracy          0.35          0.50          0.71        3331
 macro avg          0.35          0.50          0.42        3331
weighted avg          0.50          0.71          0.59        3331

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_d
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_d
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_d
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

```

Training Bi-LSTM model...
Epoch 1/5
/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/embedding.py:90: UserWarning: Argument `input_length` is deprecated. Just remove it.
  warnings.warn(
125/125 ━━━━━━━━━━━ 26s 163ms/step - accuracy: 0.7060 - loss: 0.6067 - val_accuracy: 0.6793 - val_loss: 0.5883
Epoch 2/5
125/125 ━━━━━━━━━━━ 21s 169ms/step - accuracy: 0.8018 - loss: 0.4331 - val_accuracy: 0.6358 - val_loss: 0.6193
Epoch 3/5
125/125 ━━━━━━━━━━━ 41s 168ms/step - accuracy: 0.8755 - loss: 0.3031 - val_accuracy: 0.7824 - val_loss: 0.4353
Epoch 4/5
125/125 ━━━━━━━━━━━ 39s 156ms/step - accuracy: 0.9170 - loss: 0.2196 - val_accuracy: 0.6918 - val_loss: 0.5779
Epoch 5/5
125/125 ━━━━━━━━━━━ 22s 166ms/step - accuracy: 0.9432 - loss: 0.1658 - val_accuracy: 0.7334 - val_loss: 0.5169
105/105 ━━━━━━━━━━━ 4s 29ms/step


Bi-LSTM Test Accuracy: 0.8046
precision recall f1-score support
0 0.87 0.86 0.86 2364
1 0.66 0.68 0.67 967

accuracy 0.80 3331
macro avg 0.76 0.77 0.76 3331
weighted avg 0.81 0.80 0.81 3331

```

2. Accuracy and Loss Plots

These plots illustrate training and validation performance of both models across 5 epochs. Bi-LSTM shows a stronger upward trend in accuracy and greater reduction in loss.



3. Discussion

The Bi-LSTM model did better than the LSTM model on the test data. It got over 80% accuracy, while the LSTM got about 70%. Even though their validation scores were close, Bi-LSTM had lower loss and improved more during training.

This is because Bi-LSTM can look at both the words before and after in a sentence. This helps it better understand the meaning, especially in short and unclear tweets.

Despite the LSTM being a solid baseline, Bi-LSTM is clearly more effective for this binary classification task, provided that computational resources are sufficient.